

USER MANUAL

Chhatri

छतरी · umbrella

Income cover for small Paytm merchants,
where the claim starts itself.

Build for India AI Hackathon, Track 2 (AI-powered financial journeys)

Team **Da Goats**: Omkar Kadam and Ujjwal Pardeshi

Version of 3 October 2026 · describes the prototype in this repository

Everything shown is a prototype that runs on synthetic demo data.

Contents

1. What Chhatri is	3
2. The problem, and who it is for	4
3. How it works	6
4. What is real and what is simulated	8
5. Install and run	9
6. Guided tour of every screen	12
The console (1280 by 720)	12
The merchant mini-app (390 by 844)	23
7. The three-minute demo story	31
8. Features list	34
9. Troubleshooting, FAQ and glossary	36
Troubleshooting	36
Questions a newcomer asks	37
Glossary	38

HOW TO READ THIS MANUAL


It assumes you know nothing about Chhatri. Every new word is explained the first time it appears, and there is a glossary at the end. If you only have five minutes, read section 1, then the three-minute story in section 7.

Words you will see on every page

Word	Plain meaning
Merchant	A small shop owner who takes payments on Paytm, such as a tea stall or a kirana (neighbourhood grocery).
Claim	A request to be paid because something bad happened. In Chhatri it often starts by itself.
Policy engine	The program that checks the rules and decides whether money is paid. It is ordinary code, not AI.
Simulated	Pretend but clearly labelled. For example, the WhatsApp phone is a drawing inside the console.
Demo flag	A switch that turns one feature on. All flags are off until you set them (section 5).

1 What Chhatri is

Chhatri (छहत्री, Hindi for "umbrella") is income cover for India's small merchants. When something stops a shop from earning, such as a red rain alert that empties the street, Chhatri notices the loss in the shop's own sales, pays the same day, explains the amount in Hindi, and asks the shop's lender to give that day's loan instalment a holiday. The merchant does not fill in a form. **The claim starts itself.** The AI builds the case and a rules engine written in ordinary code decides the money. Every step goes into a log that shows if anyone changed it.



The picture in one screen. It is the live map of Mumbai at 17:06 on the demo storm day, drawn on OpenStreetMap with a colour wash that shows sales against what was expected. The red zones were paid automatically: 3 zones, 312 shops, 4 minutes from trigger to money. Anil's tea stall gets ₹1,380. Sales, alerts and payouts here are simulated and labelled as such.

THE ONE-LINE VERSION

Rain stops sales. Paytm sees the drop. Chhatri pays that evening, with no form and no phone call, and the lender decides about the instalment.

2 The problem, and who it is for

The problem

A tea stall earns day by day. On a day of heavy rain nobody stops for tea, and the stall's sales can fall by more than half. The owner still has to pay his loan instalment. Illness does the same: if the owner is in hospital, the shop is shut and no money comes in.

Insurance exists for such days, but it is slow. The facts file for this project records that earlier merchant protection plans on Paytm took **30 to 60 days per claim** and needed several documents (fact A3). A merchant living on daily sales cannot wait that long.

The story used in the demo shows the gap. Anil runs a tea stall in Parel, Mumbai. His usual Tuesday sales are **₹4,380**. On the storm day his area's sales fall **63%**. His loan instalment of **₹600** is still due. With Chhatri he is paid **₹1,380** the same evening and his lender is asked to pause the instalment.

Number	What it tells you	Where it comes from
1.57 crore	Merchants who paid for Paytm payment devices in Apr-Jun 2026. These are the people Paytm already sees the sales of.	Facts file, A1
30-60 days	How long earlier merchant protection plans took per claim, with several documents.	Facts file, A3
About 250-276 mm in about 6 hours	Heavy rain in Mumbai on 25-26 Sep 2024, the kind of day that empties a street.	Facts file, A13
About 1 in 3	Chance of no payout even when the area lost everything, in a study of 270 weather-index crop insurance products. This gap between the trigger and the real loss is called basis risk.	Facts file, A8
63%, ₹4,380, ₹1,380, ₹600	Anil's demo day: sales drop, usual Tuesday, payout, instalment.	docs/DEMO.md
3 zones, 312 shops, 4 min	Demo storm result: zones triggered, shops paid, minutes from trigger to money.	docs/DEMO.md

Who uses Chhatri

The merchant

A shop owner such as Anil. Uses a phone app in Hindi or English, and gets Chhatri's messages on WhatsApp or, if he chooses, on Telegram: sees cover, gets paid, asks questions, files a complaint. Does not file claims for a storm.

The claims officer

A person who looks at the doubtful cases: a slip with the wrong name, or a merchant who says the amount is too low. Works in the console on the Claims page.

The lender


The business that gave the merchant the loan. Chhatri only *asks* it to pause an instalment. The lender decides. In this prototype the lender is simulated.

Why it matters

- **Speed.** Money arrives the same evening, not weeks later.
- **No paperwork for a storm.** An area claim needs no document at all. A hospital claim needs one photo of the hospital slip.
- **Fairness you can check.** Every amount is a sum the merchant can read: half of the usual day, times the drop.
- **Humans stay in charge of doubt.** If the AI is not sure, a person decides.

3 How it works

Follow one claim from left to right. A **trigger** (something that starts a claim) goes to the **policy engine**, which reads a rules file called `rules.yaml`. If the rules say yes, money is paid and the lender is asked about the instalment. Everything is written to the **audit log**.



The two kinds of claim

	Area claim (storm)	Personal claim (illness)
What starts it	A red or orange weather alert, and the hourly sales index of a zone stays below 50% of expected for 3 hours in a row, with at least 20 shops in the zone.	A shop shows a full day with no sales. Chhatri checks in by chat in Hindi (on WhatsApp, or on Telegram if the merchant chose it). The merchant sends one photo of the hospital slip.
Documents	None.	One hospital slip.
Amount	Half of the lost sales, capped at ₹2,500 per shop per day.	Half of the usual day, capped at ₹1,500 a day, up to 3 days paid automatically.
Example	$\frac{1}{2} \times ₹4,380 \times 63\% = ₹1,380$ (Anil).	$\frac{1}{2} \times ₹4,300 = ₹2,150$, capped at ₹1,500 for 1 day = ₹1,500.

THE KEY WORDS

Expected sales: what the shop normally sells at this hour, estimated by a small statistical model. **Sales index:** actual sales divided by expected sales. **Zone:** a group of neighbouring shops in one area; the demo uses Z3, Z7, Z9 and Z12. **Settlement:** the daily payment Paytm makes to a merchant. **EDI:** the equated daily instalment of a loan, taken out of the settlement. **EDI holiday:** a pause of one instalment, which moves to the end of the loan.

Who is allowed to decide what

Case	Chhatri alone	Goes to a human
Area drop during an alert, index clear	Pays	Only if the merchant disputes
Personal claim, slip matches name and dates	Pays up to the daily cap	Anything above the cap

Case	Chhatri alone	Goes to a human
Slip unclear or dates do not match	Never	Always
Cover bought after an alert	Never	A waiting period applies

This is the table on the Policy page (screen 6.15). Other rules in `rules.yaml`, version `pilot-0.1`: payout share 50%; slip name must match the KYC name (score 85 or more); slip reading confidence 0.80 or more; waiting period 7 days; annual limit ₹30,000; disputes answered within 24 hours. **KYC** means the identity record of the merchant.

Where AI helps, and where it never decides money

AI helps with	How	It never
Reading a hospital slip	A vision model reads the name, dates and hospital. The merchant confirms what was read.	Approves a claim. The engine checks the name against KYC.
Ask Chhatri	Answers questions in plain words and cites the policy clause. Known questions are answered by rules; free questions go to Gemini, then Sarvam, then a fixed template. A guard blocks any model reply with an invented amount or a promise. Every number comes from the engine.	Invents an amount or promises a payout.
Voice	Listens and speaks in Hindi and English. Amounts and dates are confirmed with a tap.	Acts on an amount the merchant has not confirmed.
Hidden instructions	Text found inside slips and chat is treated as data, never as an order.	Obeys it.

THE RULE THAT MATTERS

The AI builds the case. Code decides the money. A message that tells a merchant an amount is filled from a fixed template with numbers from the engine, never written freely by a model.

4 What is real and what is simulated

Chhatri is a prototype, and it says so on screen. Every outside service has a small label in the console header (click the chip that reads *Simulated · 18*; a live service is named there instead). There are three labels:

Label	Meaning
LIVE	A real service is being called (an AI key or the Telegram token is set, or n8n runs under <code>make up</code>).
SIMULATED	A simulator answers instead. Nothing outside is called.
FALLBACK	The service was meant to be live but failed or was switched off, so a simpler answer is used. This label appears with the provider panel flag.

Part	Status in this prototype
Policy engine, forecast model, audit log, claims console	Real code, running for real (on simulated shop data).
Shop sales, weather alerts feed, KYC, payout rail, lender, Soundbox (the speaker on the till)	Always simulated and always labelled.
WhatsApp and Telegram	Simulated. The phone you see is drawn inside the console. WhatsApp becomes live only with four WhatsApp settings, which this project has not run. Telegram becomes live only with <code>TELEGRAM_BOT_TOKEN</code> and the flag <code>telegram_channel</code> , for demo merchants.
Paytm payment link	Simulated. A pretend link such as <code>https://paytm.me/sim-...</code> is shown and labelled.
Sarvam (speech, chat, slip reading)	Simulated unless <code>SARVAM_API_KEY</code> is set.
Gemini (chat and slip reading)	Simulated unless <code>GOOGLE_API_KEY</code> and <code>GEMINI_MODEL</code> are set.
Workflows (n8n)	An in-process runner does the same steps. The n8n label reads simulated unless you run the Docker stack.
Weather	Rain data in the backtest is real, from Open-Meteo, saved in the repository. The live weather call is off by default.
Map pictures	Real OpenStreetMap tiles from the internet, with simulated sales drawn over them as colours. Offline, the console shows plain ward outlines.

HONEST LIMITS

- AI services and the Telegram bot run **only on synthetic demo data**. Without `CHHATRI_DATA_IS_SYNTHETIC=true`, every free-tier AI call is skipped.
- The backtest (section 6.13) uses **simulated sales with real rainfall**. It tests the rule, not accuracy on real shops. Its simulation was tuned so the demo numbers appear, and premiums are priced so each zone ends near a 65% loss ratio, so those numbers are by design, not a finding. The pricing simulator (6.14) gives planning figures, not an actuarial price.
- No insurer, lender or Paytm team has agreed to take part. It is not an approved insurance product.
- AI accuracy is only partly measured: the one stored run (3 October 2026) shows two misses, slip reading and voice are NOT MEASURED, and no person has graded an answer (section 6.16).

5 Install and run

What you need first

- **Python 3.12** and **Node 22** (Vite, the web tool, needs Node 20.19 or newer) and **GNU make**.
- **Docker with Compose v2**, only if you use `make up`.
- No API key is needed. Without keys every service is simulated.

Run everything on your machine

```
git clone <repo> chhatri && cd chhatri
make setup      # backend/.venv with the backend, and npm ci in frontend/
make env        # creates .env with two generated secrets (never overwrites an existing one)
make dev        # backend on http://localhost:8000, console on http://localhost:5173 (Ctrl+C stops both)
```

Open `http://localhost:5173`. The backend loads the **monsoon** scenario at start, paused at 08:00. `GET http://localhost:8000/api/preflight` lists what is ready. `make data` rebuilds the saved model files; it is slow and needed only if a file is missing, because the results are committed.

Run with no backend at all (the static mock)

The console can run a copy of the whole backend inside the browser. Nothing else needs to run, so it works offline and can be put on a free static host.

```
cd frontend
npm ci
npm run dev:mock                # development server in mock mode

# a static build, shown exactly with the flags it was built with:
VITE_FEATURES=n1_miniapp npm run build -- --mode mock
npm run preview                 # then open /merchant/S-0142/app
```

In mock mode every service is simulated and the footer shows a "Mock data" badge. A full page reload restarts the mock, so move around with the links, not the reload button.

Check that it works

Command	What it does
<code>make judge</code>	One command to check the whole build. Six checks, each printing PASS, FAIL or INFO: the environment, the keys (SET or NOT SET is never a failure), the saved data files against their recorded fingerprints, the demo check (70 checks), the audit chain after the payouts, and the frontend type check. It ends with <code>CHHATRI JUDGE READY ✓</code> when nothing failed, takes about a minute and calls no outside service.
<code>make demo-check</code>	Runs every demo scenario through the HTTP API and prints <code>PASS: 70 passed, 0 failed, 0 skipped</code> when all is well. It runs on its own copy of the app with only <code>x4_lender_request</code> on; do not turn on <code>n3_slip_precheck</code> or <code>n6_consent</code> s for it, because its scripted steps stop by design when those two are on.

Command	What it does
<code>make check-keys</code>	Prints whether <code>SARVAM_API_KEY</code> , <code>GOOGLE_API_KEY</code> and <code>TELEGRAM_BOT_TOKEN</code> are SET or NOT SET, and never a key. With a Google key it also lists the Gemini models the key can use.
<code>make test</code>	Backend tests (coverage gate 80%) and frontend type check, lint and unit tests.
<code>make e2e</code>	Playwright browser tests against running servers.
<code>make evals</code>	Runs the offline AI evaluation suites (no network, no key) and writes a new summary for the Evals page. <code>python -m chhatrival -m chhatrival --live</code> also runs the 50 Ask questions against a live model.
<code>make demo-stage</code> <code>make stage-e2e</code>	<code>demo-stage</code> starts the stage demo of section 7: backend and console together, the stage flag set, synthetic data, no auto-reload (<code>STAGE_AI=sim</code> blanks both AI keys). <code>stage-e2e</code> walks the same script with a browser test on fresh servers (ports 8301 and 5301), by default with no AI keys, then stops them.
<code>make up</code> <code>make down</code>	Docker stack: backend, console on http://localhost:8080 and n8n on http://localhost:5678 .

API keys (optional)

Put keys in `.env` (the file `make env` creates). Leave them out and everything stays simulated. `make check-keys` shows which keys are set.

Setting	Turns on
<code>GOOGLE_API_KEY</code> and <code>GEMINI_MODEL</code>	Gemini, the first AI link for Ask Chhatrival and slip reading. A key alone is not enough: a model id is needed too (<code>GEMINI_VISION_MODEL</code> is optional; <code>make check-keys</code> lists the models your key can call).
<code>GEMINI_BACKUP_MODELS</code>	Optional list of more Gemini model ids, separated by commas. The free tier allows only about 20 requests a day for each model, so when one is used up the next answers before Sarvam or a fixed template does.
<code>SARVAM_API_KEY</code>	Sarvam speech to text, text to speech, chat and slip reading. It is the next link after Gemini, and the one that gives Hindi voice.
<code>CHHATRI_DATA_IS_SYNTHETIC</code>	Set to <code>true</code> : the declaration that all data is synthetic (already in <code>.env.example</code>). Free-tier AI (Gemini, Sarvam) and the Telegram bot get synthetic data only: with this line off, every such call is skipped and the label reads SIMULATED. Never set it where real merchants' or patients' data is held.
<code>TELEGRAM_BOT_TOKEN</code>	A live Telegram bot (flag <code>telegram_channel</code>). Make a bot with Telegram's @BotFather and paste its token; the bot reads messages by long polling, so no public address is needed. To link a demo merchant, open <a href="https://t.me/<bot_name>?start=S-0142">https://t.me/<bot_name>?start=S-0142 on a phone and press Start. Only one machine may poll a token: set <code>TELEGRAM_POLLING=false</code> on the others.
The WhatsApp, Paytm and n8n settings	Live WhatsApp (four settings), a live Paytm staging link (<code>PAYTM_MCP_URL</code> , or <code>PAYTM_MID</code> and <code>PAYTM_KEY_SECRET</code>), n8n workflows (<code>N8N_BASE_URL</code>). Not used in the demo.

How the AI links are tried: for a free question the rules answer first, then Gemini, then Sarvam, then a fixed template; for a slip, Gemini vision, then Sarvam vision, then a person. The quota above was about 20 requests a day for each model when this was written (3 October 2026). Rehearsals spend the same quota as the show, so do not loop them on a live model.

The 15 feature flags

A **feature flag** is a switch. Every feature added for the final is off until you name it. Off means absent: the screen is hidden and the backend answers an ordinary 404. Set the same list in two places:

`CHHATRI_FEATURES` for the backend (in `.env`) and `VITE_FEATURES` for the console (in the shell or in `frontend/.env.local`), separated by commas.

```
CHHATRI_FEATURES=n1_miniapp,n2_ask_chhatri,n3_slip_precheck,n4_voice,n5_grievances,n6_consent,  
n8_marathi,x4_lender_request,x6_provider_panel,x8_distress_guard,h8_ops_strip,h24_whatif,h25_evals,  
console_polish,telegram_channel  
# the console needs the same names:  
VITE_FEATURES=<the same list> npm run dev:mock
```

Flag	What it turns on
<code>n1_miniapp</code>	The merchant mini-app: home, coverage, buy cover, claims tracker, claim detail, why this amount, receipt, help, language switch (Hindi, English), jargon lens, next-step bar.
<code>n2_ask_chhatri</code>	Ask Chhatri: questions answered with policy clause citations and a scam-message warning.
<code>n3_slip_precheck</code>	The slip pre-check: the app shows what it read and a checklist, and the merchant confirms before any check runs.
<code>n4_voice</code>	Voice: the microphone and the confirm chips for amounts and dates.
<code>n5_grievances</code>	The complaint ladder with response clocks.
<code>n6_consent</code>	The consent centre, the consent activity log, "Erase this slip", and the consent tick boxes when buying cover.
<code>n8_marathi</code>	Marathi as a third language (a draft, text only).
<code>x4_lender_request</code>	The lender decides the EDI holiday (without it the 17:05 line is an unconditional pause). The demo flag.
<code>x6_provider_panel</code>	The provider panel with Force fallback switches and the FALLBACK label.
<code>x8_distress_guard</code>	No loan offers to a merchant in distress, and a cap on how many messages one merchant gets.
<code>h8_ops_strip</code>	The strip of five real operations numbers at the top of the console.
<code>h24_whatif</code>	The "What if..." drawer on a zone card, and the pricing simulator "Price the cover" at the end of the Backtest page.
<code>h25_evals</code>	The Evals page (AI evaluation).
<code>console_polish</code>	Presenter mode, the trigger-to-payout card and projector-sized type.
<code>telegram_channel</code>	Telegram as a second chat app next to WhatsApp, which stays the default: the Chat app switch on the Merchant phone page, "Messages from Chhatri" in the mini-app's Settings, and a Telegram row in the provider panel. A live bot needs <code>TELEGRAM_BOT_TOKEN</code> (see API keys above).

You can see which flags are on with `GET /api/health`. Run `make dev` or `npm run dev:mock` with all 15 flags on to see every screen in section 6; `make demo-stage` turns on only the flags the three-minute story needs.

6 Guided tour of every screen


These screens were captured from the static mock with all 15 flags on, except 6.14, which needs the backend's pricing table and was captured from the full stack. The console pictures are 1280 by 720 pixels (a projector). The phone pictures are 390 by 844 (a phone). Each picture says what it shows and what to do. Times are simulated clock times of the demo day, Tuesday 19 August 2025.

The console pages are: Overview, Live map, Claims, Merchant phone, Audit, Backtest, Policy, Evals (the menu across the top). The small clock bar under the menu moves the demo day. The footer of every console page names the data sources (OpenStreetMap map, DataMeet wards, Open-Meteo rainfall), says which parts are simulated, and shows a *Mock data* badge while the static mock is running.

The console (1280 by 720)

6.1 Overview

Menu: **Overview** · the front page



The front page, built as a short product page. It opens with the promise, *The claim starts itself*, three buttons and a phone with the chat Anil sees. A strip of four figures from the storm replay follows (same day, 4 minutes to the credit, 312 shops paid, 0 claim forms), captioned as real rainfall with simulated sales. Scrolling down you meet the problem, how it works (Detect, Decide, Pay, Protect), who decides (second picture: AI understands, rules decide, Paytm executes), why Paytm, the storm example, the merchant phones, the product's features, the backtest proof, who gains and what it costs, what is real and what is simulated, and common questions.

1. Press **Watch the storm replay**. It loads the monsoon day, plays the last minutes before the trigger on the Live map and pauses at 17:06.
2. Press **See Anil's phone** to open the Merchant phone page at 17:05 (6.11).
3. Press **Open the claims console** to load the slip-mismatch case in the officer queue (6.10).
4. In the header, **Simulated · 18** opens the provider panel (6.9), **Enable sound** lets the Soundbox voice play (click it once) and **Present** turns on Presenter mode (6.8).

6.1 Overview, scrolled down: Who decides

Chhatrī Overview Live map Claims Merchant phone Audit Backtest Policy Evals Present Simulated 18 Enable sound

WHO DECIDES

AI understands. Rules decide. Paytm executes.

AI is allowed to be unsure. It is never allowed to move money: a payout needs the policy engine, and an unclear case goes to a person.

AI UNDERSTANDS
Language and documents

- ✓ Hears Hindi and Marathi voice notes
- ✓ Reads the hospital slip
- ✓ Answers "why this amount?" from the policy, with clause numbers
- ✓ Spots scam and prompt-injection messages

Gemini or Sarvam when keys are set; written templates when not.

RULES DECIDE
The policy engine

- ✓ The only part that can say APPROVED
- ✓ Same facts in, same decision out
- ✓ Lists every check and clause it used
- ✓ Sends anything uncertain to a claims officer

Deterministic code with a published rule set (pilot: 0.1).

PAYTM EXECUTES
Money and messages

- ✓ Credits the payout with the settlement
- ✓ Announces it on the Soundbox
- ✓ Asks the lender to pause the instalment
- ✓ Tells the merchant on WhatsApp or Telegram

Simulated in the prototype; the Telegram bot can run live.

Map: © OpenStreetMap contributors (ODBL) Wards: DataHeet (CC BY-SA 2.5 India) Rainfall: Open-Meteo Sales, alerts, KYC, payouts, lender and Soundbox are simulated **MOCK DATA**

6.2 Live map, loaded at 08:00

Menu: Live map

Chhatrī Overview Live map Claims Merchant phone Audit Backtest Policy Evals Present Simulated 18 Enable sound

Monsoon replay - Tue 19 Aug Mumbai: monsoon replay 08:00 simulated Play 6 min/s +1 min +15 min +1 h HHMM Seek Slow near payout

Open cases 0 none open Next due Nothing due Decided by the engine No claims yet Paid today ₹0 · 0 shops Holiday requests 0 granted

Red alert from 14:00 - 3 zones

Zone 7 - 46 shops

100% of expected sales

Alert Red alert from 14:00

Sales 100% of expected for 3 hours

Cover 46 of 46 prepaid

0 zones triggered 0 shops paid trigger to money

Z9 - slow day, no alert

Live events streaming

08:00 Z9 at 64% with no alert: slow day, no payout

08:00 Red alert A-20250818-01 for Z3, Z7, Z12 - valid 14:00-20:00 (issued Mon 18 Aug 17:30)


Map: © OpenStreetMap contributors (ODBL) Wards: DataHeet (CC BY-SA 2.5 India) Rainfall: Open-Meteo Sales, alerts, KYC, payouts, lender and Soundbox are simulated **MOCK DATA**

Mumbai before the storm, on an OpenStreetMap map so the real place names stay readable. The colour laid over it shows each zone's sales against what was expected: green is normal, amber is falling, red is below the 50% line (the key is at the bottom right). Before any rain the map opens wide on the whole Mumbai Metropolitan Region (MMR); the faint colour outside the city is simulated context for areas Chhatrī does not cover. The card on the right is Zone 7, *Normal*. The bar at the top is the demo clock.

1. Pick a day in the **scenario** list at the top left. Choose **Monsoon replay**. The clock reads *Mumbai · monsoon replay · 08:00 · simulated*.
2. Type a time such as 15:30 in the box and press **Seek** to jump there. Press **Play** to run it; the speed list sets how fast (6 min/s means 6 simulated minutes per second). **+1 min**, **+15 min** and **+1 h** step the clock; the round arrow resets.
3. The marks under the bar are the story: 14:00 alert, 17:00 trigger, 17:04 paid, 17:05 instalment. **Slow near payout** (on by default) plays 16:58 to 17:06 at one simulated minute per second so you can watch.
4. The switch at the top right of the map flips between **Mumbai** (the storm wards) and **MMR** (the whole region). Once rain appears the map moves in on Mumbai by itself; your own choice always wins.

6.3 The storm builds

Live map · seek to 15:30




The red rain alert has been on since 14:00 for zones Z3, Z7 and Z12, and a rain band is drawn over them (the hatched blue, with the pill *Heavy rain band · since 14:00*). Their sales are falling, so the colour moves from green through amber towards red, and the map has moved in on Mumbai. Zone Z7 shows *Watching*.

1. Seek to **15:30**.
2. Look at the three zone labels on the map. A label such as *Z7 · 66% · 46 shops* reads: sales are 66% of expected right now, 46 shops. The card on the right shows the same zone over the last three full hours (76% for 12:00 to 15:00) and the live hour (*Now: 66% (live)*).
3. Look at Z9 on the east side. It dips too, but there is no alert there.

6.4 The trigger fires, and the zone panel

Live map · seek to 17:06



At 17:00 the drop has held below 50% for 3 hours during the alert, so the trigger fires for Z3, Z7 and Z12 and their zones turn red. The three tiles on the right read 3 zones triggered, 312 shops paid, 4 minutes from trigger to money. A navy note at the top of the map says **₹1,380 credited · 17:04** for Anil's Tea Stall, and his pin reads **₹1,380 paid · 17:04**.

1. Seek to **17:06** (or park at 16:40, press Play and wait about 12 seconds).

2. The card on the right is the **zone panel** for Zone 7: alert, sales (37% of expected for 3 hours), cover (46 of 46 prepaid), paid (17:04, with the settlement), total ₹58,900, instalments paused. Click another zone label to change the card.
3. Read the blue note at the bottom: *Why Zone 9 got nothing*. Its sales fell to 61% with no alert. A slow day is not a loss event, so Chhatri does not pay.
4. The strip above the map is the **ops strip** (6.5). The bar above the zone panel is the **moment card** (6.6).

6.5 The ops strip

Top of the Live map and Claims · flag `h8_ops_strip`


Open cases 0 none open	Next due Nothing due	Decided by the engine 100% · 312 of 312	Paid today ₹0 · 0 shops 312 in flight	Holiday requests 0 granted
----------------------------------	--------------------------------	---	---	--------------------------------------

Five real counts, taken from the demo database: **Open cases**, **Next due** (the case whose 24-hour answer is nearest), **Decided by the engine**, **Paid today** and **Holiday requests**. This picture is from 17:01, when 312 payouts are still "in flight". Once paid, Paid today reads ₹4,25,420 for 312 shops (the Indian way of grouping digits).

1. Seek to 17:01 and watch **Paid today** read ₹0 with "312 in flight".
2. Seek to 17:06 and see it fill in.
3. Click a tile for a small pop-up with the detail.

6.6 The moment card

Right of the Live map · flag `console_polish`




A small timeline from trigger to money. At 17:01 it reads *Triggered at 17:00 · paying 312 shops · credit due 17:04*. Once paid it turns green: *123 holiday requests, 123 granted*.

1. Seek to 17:01, then press **+1 min** four times and watch the dot move along Trigger, Paid, Instalment.

6.7 The what-if drawer

Zone card · button **What if...** · flag `h24_whatif`




A read-only calculator for judges. It re-runs the real trigger rule for Zone 7 with numbers you change. The pill *Read-only: nothing is saved* says it. Nothing is saved and nobody is paid.

1. Click **What if...** on the zone card. The drawer says when it was evaluated (*Evaluated at 17:00, window 14:00 to 17:00*); **Latest hour** moves it to the newest hour.
2. Choose the alert type (None, Rain, Civic or Heat) and drag the three hour sliders (sales as a percent of expected; the red mark on each is the 50% floor).
3. Change **Shops in the index** (the rule needs at least 20) or tick **Already triggered today**.
4. Read the table at the bottom, which shows each condition for *What happened* and *What if*. Press **Back to what happened** or **Close** to leave.

6.8 Presenter mode

Header button **Present** · flag `console_polish`




A bigger-type view for a projector. The type grows, **Play** stays in view, the speed, step, seek and reset controls fold behind a **...** button, and a **Keys** button shows the keyboard shortcuts.

1. Press **Present** in the header. It stays pressed while the mode is on.
2. Press **Keys** for the keyboard shortcuts, and **...** for the speed list, step buttons, seek box and reset.
3. Press **Present** again to go back to the normal size.

6.9 Provider panel

Header chip **Simulated** (or Live) · flag `x6_provider_panel`




The list of every outside service and whether it is live, simulated or fallback: 18 rows with the Telegram flag on, from Sarvam speech and WhatsApp to Gemini and Telegram. It is how the project proves that nothing simulated is shown as live. The top line says *Static demo: nothing is live, so no row can be forced*. When a key is set, the header chip names what is live instead (for example *Gemini · Telegram + 15 simulated*, as in the picture of 6.14).

1. Click the chip in the header that reads **Simulated · 18**.
2. Read each row: name, label, a one-line reason (for example *Slip reader reads the sample slips' embedded data*) and a short reason code.
3. With a real backend the **Force fallback** button switches a live service off, so you can show the fallback on purpose. A forced service is marked *forced*; **Clear all** undoes it. Press **Esc** to close.

6.10 Claims queue and the case panel

Menu: **Claims** (a red number counts open cases)




Where the claims officer works. On the left is the queue. On the right is the open case: here **C-2291**, a dispute. Anil said "my loss was bigger". The panel shows the engine's decision (₹1,380, the formula, the decision number), a line saying one more point of drop would have added about ₹22, the lender's answer on the instalment, and the evidence chart.

1. To get a case: open **Merchant phone**, tap the Hindi voice chip that means *my loss was bigger*, then follow the blue chip **Sent to a claims officer · case C-2291**.
2. Read the grey decision box. Under it the page says **The amount cannot change**.
3. Type an optional note, then press **Confirm payout** (keeps the payout) or **Reject dispute** (closes it). The merchant is told either way.
4. For a REFERRED slip case (a slip with the wrong name) the officer presses **Approve**; the approval re-runs every hard check and then pays ₹1,500.

6.11 Merchant phone and the chat app

Menu: **Merchant phone** · seek to 17:06 · flag telegram_channel for the Chat app switch



Anil's side of the story on one page. On the left is the chat as it looks on WhatsApp (the payout card, the lender line, voice chips). In the middle is the mini-app (6.17 onwards). On the right are the facts behind it: the Soundbox line, the money with its formula, and what happened at 17:00, 17:04 and 17:05. The **Chat app** switch chooses where Chhatrri's messages go: **WhatsApp**, the default, or **Telegram**. The second picture is the page after switching to Telegram: the phone turns blue and is named *Chhatrri bot*, and the line under the switch reads *Telegram · SIMULATED · no chat linked yet*. Only where messages go changes: cover and payouts stay the same.

1. Tap a voice chip on the phone, for example the Hindi chip that means *Why did I get only this much?*. The reply shows the formula.
2. Tap the chip that means *my loss was bigger* to open a case for a claims officer (6.10).
3. Press **Telegram** in the **Chat app** switch (**WhatsApp** switches back). In the static demo the bot is simulated, so its messages show on this phone; with a live bot a link, *Open @bot name*, opens it on a real phone. Loading a scenario starts again on WhatsApp.

6.11 Merchant phone, after pressing Telegram in the Chat app switch

The screenshot displays the Chhatri Merchant phone interface. At the top, there's a navigation bar with tabs: Overview, Live map, Claims, Merchant phone (active), Audit, Backtest, Policy, and Evals. Below this is a timeline bar showing a simulated replay of a monsoon event on Tuesday, August 19th, in Mumbai. The timeline includes markers for 14:00 Alert, 17:00 Trigger, 17:04 Paid, and 17:05 Instalment. The main area is divided into three panels. The left panel shows a Telegram chat with 'Chhatri bot' where a claim officer has been contacted. The middle panel shows a simulated Telegram chat with a merchant named 'नमस्ते, अनिल जी' (Namaste, Anil Ji) discussing a payment of ₹1,380. The right panel shows a simulated payment confirmation screen with details like 'Paytm par ₹1,380 prapt hue' and 'Chhatri se'.

6.12 Audit log with Verify chain

Menu: Audit

The screenshot displays the Chhatri Audit log interface. At the top, there's a navigation bar with tabs: Overview, Live map, Claims, Merchant phone, Audit (active), Backtest, Policy, and Evals. Below this is a timeline bar showing a simulated replay of a monsoon event on Tuesday, August 19th, in Mumbai. The timeline includes markers for 14:00 Alert, 17:00 Trigger, 17:04 Paid, and 17:05 Instalment. The main area is titled 'Audit log' and contains a search bar with the text 'Search: payout, lender, S-0142...'. Below the search bar is a 'Verify chain' button and a status message: 'Chain valid - 22 entries - none changed since written'. The main content is a table with columns: #, Who, What happened, Record, Fingerprint, and Previous. The table lists 22 entries, grouped by simulated minute (19 Aug 2025 - 17:06 and 19 Aug 2025 - 17:05). The entries include messages sent, case opening, and payment workflow steps.


#	Who	What happened	Record	Fingerprint	Previous
19 Aug 2025 - 17:06					
22	Chhatri assistant	Message sent	Message M-000007	27a12c0c93	c4d26daa30
21	Chhatri assistant	Message sent	Message M-000006	c4d26daa30	8c794b76eb
20	Chhatri system	Opened a case for a claims officer dispute	Case C-2291	8c794b76eb	fd4c152004
19	Chhatri assistant	Understood a message from the merchant dispute amount	Merchant S-0142	fd4c152004	bd68a3b008
18	Merchant S-0142	Message received	Message M-000005	bd68a3b008	2173aa2e23
19 Aug 2025 - 17:05					
17	Chhatri assistant	Message sent	Message M-000004	2173aa2e23	5587eb1f42
16	Payout workflow	Paused the ₹600 instalment of the merchant	Instalment pause IP-000001	5587eb1f42	fa84f45be7

A tamper-evident log in plain words and simulated time. Each entry stores a fingerprint (a sha256 hash) of the one before it, so changing any entry breaks the chain. The newest entry is at the top, grouped by simulated minute; each row says who acted, what happened, the record it is about, its fingerprint and the one before it.

1. Press **Verify chain**.
2. A green note reads *Chain valid · 22 entries · none changed since written* (your count may differ) and every row gets a green tick.
3. Type in the search box to find a payout, a lender or a merchant, for example `payout` , `lender` or `S-0142` .
4. Read the rows from the bottom up (the oldest entry is at the bottom): trigger, decision, payout, holiday request, lender answer, message.

6.13 Backtest

Menu: Backtest




Would Chhatri have paid the real losses of the 2024 and 2025 monsoons? It compares Chhatri with a weather-only trigger. It is labelled *simulated sales · real Open-Meteo rainfall*.

1. Left chart: real sales drops that get paid. Chhatri paid 89 of 148 (60%); a weather-only trigger paid 49 (33%).
2. Right chart: payouts with no real drop. Chhatri 36 of 125 (29%); weather-only 287 of 336 (85%).
3. Below: trigger to money (same day, 4 minutes), documents per area claim (0), doubtful personal claims seen by a human (all of them: 359 of 844 personal claims went to a human, and 485 clean claims were paid automatically).
4. Scroll down for every measure, premiums against payouts per zone and, at the very end, **Price the cover** (6.14). Read the line under the chart: the model tests the rule, not accuracy on real shops.

6.14 Price the cover

Menu: Backtest, last section (/backtest#pricing) · flag h24_whatif



A read-only price calculator for the question "what if the price is wrong?". Every zone is priced from the backtest's own triggers (simulated sales, real rainfall). It opens on the published rules (pilot-0.1) for Zone 7, Parel · Lalbaug: index floor 50%, payout share 50%, area daily cap ₹2,500, loading 35%. The tiles give the premium (₹18.62 a day, the same price Anil sees on Home), the expected payout per shop, the payout days, the covered shops and the loss ratio at today's price (65%). Under them, across the city the premium runs

from ₹6.93 (lowest) through ₹16.43 (median) to ₹38.82 (highest), and at a 50% floor 89 of 148 real drops are paid while 36 of 125 payouts fall on a day with no real drop. This picture comes from the full stack; the static mock only says the simulator needs the backend's pricing table.

1. Scroll to the end of the Backtest page, or follow **Try the pricing simulator** from the Overview.
2. Press another **Index floor**, or drag **Payout share**, **Area daily cap** or **Loading**; a lever that differs from the published rules says what the rules have (*published 50%*). Pick another zone in the **Zone** list. A loss ratio above 100% means the zone pays out more than it collects.
3. Press **Reset to the published rules** to go back. Nothing is saved, and the caption says *Not an actuarial price*: the real price is the open question for a pilot.

6.15 Policy

Menu: Policy

The screenshot shows the Chhatrī Policy page. At the top, there's a navigation bar with tabs: Overview, Live map, Claims, Merchant phone, Audit, Backtest, Policy (selected), and Evals. Below the navigation bar, there's a simulation timeline for 'Monsoon replay - Tue 19 Aug' in 'Mumbai'. The timeline shows a '14:00 Alert' and a '17:00 Trigger'. The current time is '17:06' and the simulation is 'simulated'. There are controls for 'Play', '6 min/s', '+1 min', '+15 min', '+1 h', '17:06', 'Seek', and 'Slow near payout'.

The main content area is titled 'Automatic when the data is clear, human when it isn't'. It includes a 'Payout authority' table, 'Checks before any payout', and 'Three live tests in our demo'.

Case	Chhatrī alone	Goes to a human
Area drop during an alert, index clear	Pays	Only if the merchant disputes
Personal claim, slip matches name and dates	Pays up to the daily cap	Anything above the cap
Slip unclear or dates don't match	Never	Always
Cover bought after an alert	Never	Waiting period applies

New cover starts only after a **waiting period**, so nobody can buy it once a storm is forecast.

Checks before any payout

Check	Applies	Severity	Passes when
COVER IN FORCE	all	HARD	cover exists, state on or before the event

Map: © OpenStreetMap contributors (ODBL) Wards: DataMeet (CC BY-SA 2.5 India) Rainfall: Open-Meteo Sales, alerts, KYC, payouts, lender and Soundbox are simulated **MOCK DATA**

Three live tests in our demo

"My loss was bigger than that."
EXPLAINED Chhatrī shows the numbers and offers a human review.
Run it →

A hospital slip with a different name
HUMAN No automatic payout. A claims officer decides.
Run it →


"Red alert tomorrow. Cover me today."
BLOCKED New cover starts after the waiting period.
Run it →

The page headed *Automatic when the data is clear, human when it isn't*. It holds the payout-authority table (who may decide what) and the checks that run before any payout. On the right are the three live tests of the demo: EXPLAINED, HUMAN and BLOCKED.

1. Read the table: green means Chhatrī pays alone, red means never alone.
2. Press **Run it** on a test card to load that scenario.
3. Scroll to **Checks before any payout**: each check, where it applies, whether it is HARD, and when it passes.


6.16 Evals

Menu: Evals · flag h25_evals





The AI evaluation page, headed *Measured, not claimed*. It shows one stored run (`run-20261003T080226` , 3 October 2026) as six suites: S1 intent routing, S2 guard red-team, S3 Ask end to end, S4 slip reading and the gate, S5 voice, S6 chains and labels. Every line is a count (for example 51 of 56 messages routed as labelled) with a 95% interval, a target where there is one, and a label: MEASURED, MET, *MET*, *WIDE INTERVAL*, MISSED or NOT MEASURED (a suite with a part not run, such as S3, reads PARTIAL). The data is synthetic, as the note with the amber edge says.

1. Read the run line under the title: the run id, the code version, the times, and what ran as rules (S1, S2), live (the Ask suite ran through the live chain on gemini-flash-lite-latest; 11 of 48 questions still ended in a template) or as a mock.
2. Read the misses, which the page shows plainly: in S2, 1 of 35 replies that had to be blocked got through; in S3, 2 of 50 Ask answers held a rupee figure that no fact supported. Both read MISSED.
3. Scroll down (second picture): answers rated by two reviewers are NOT MEASURED because no person has graded an answer yet, and S4 slip reading and S5 voice are NOT MEASURED because they need a live provider run. A new run replaces the stored one (`make evals` runs the offline suites).



The merchant mini-app (390 by 844)

The mini-app is the merchant's side of Chhattri: a phone-sized app inside the console (page **Merchant phone**) and also full screen at `/merchant/S-0142/app`. The first merchant is Anil (S-0142), shown at 17:05 on the storm day. Three tabs at the bottom (Home, Claims, Help) move between the main areas; the bar above them always suggests the next step, and its words change with the situation. The top bar shows the name of the screen (on Home, the brand *Chhattri*), the demo date and time, a SIMULATED label and the globe icon, which opens the language and chat app settings. On the static demo a blue banner says *This is the static demo*; press **Got it** to fold it away (it is folded away in the pictures below). The pictures are in English, except where Hindi is named.



6.17 Home

Tab Home

Anil's front screen. The top bar names the app (*Chhattri*) and shows the demo date and time. Below it: the weather alert for his area and his cover (active, paid through 22 August, ₹18.62 a day, ₹600 loan instalment, used ₹1,380 of ₹30,000). The right picture is the screen scrolled: his latest claim and two shortcuts.

1. Read the amber alert box: you do not need to do anything.
2. Read the dark cover card.
3. Press the blue bar, **Why this amount?**, to see how ₹1,380 was worked out.

What am I covered for?
Demo date and time: 19 August, 17:05

◀ SIMULATED ⓘ

Your cover in numbers

Chhatri's share of lost sales	50%
Daily limit, rain	₹2,500
Daily limit, hospital cash	₹1,500
Yearly limit	₹30,000
Waiting period	7 days
Alerts checked this far ahead	72 hours
First payment covers	30 days

Tap an underlined word to see what it means.

[Waiting period](#) [Chhatri's share](#) [Daily limit](#)

[Yearly limit](#) [Weather alert](#)

Rain claims need no forms. Hospital cash needs one photo of the slip.

Rain and lost sales C2

When a weather alert is on, your area's sales stay below 50% of the usual level for 3 hours in a row, and the fall is bigger than an ordinary slow day.

See what Chhatri has done for you.

[See my claims](#)

Home Claims Help

What am I covered for?
Demo date and time: 19 August, 17:05

◀ SIMULATED ⓘ

Your cover in numbers

Chhatri's share of lost sales	50%
Daily limit, rain	₹2,500
Daily limit, hospital cash	₹1,500
Yearly limit	₹30,000
Waiting period	7 days
Alerts checked this far ahead	72 hours
First payment covers	30 days

Tap an underlined word to see what it means.

[Waiting period](#) [Chhatri's share](#) [Daily limit](#)

Waiting period
वेटिंग पीरियड

In plain words
A new cover starts 7 days after you ask. Until then it does not pay.

Example
Ramesh asks on Monday 18 August. His cover starts on 25 August.

This explains the word in plain language. The policy wording has the exact terms.

[Close](#)

6.18 Coverage, and the jargon lens

Home · What am I covered for?

The cover in numbers: Chhatri pays 50% of lost sales, up to ₹2,500 a day for rain and ₹1,500 a day for hospital cash, ₹30,000 a year, a 7-day waiting period, 72 hours of alert look-ahead and a first payment that covers 30 days. Underlined words are explained on tap: the **jargon lens**, a sheet with plain words, an example and the other language.

1. Scroll through **Your cover in numbers**.
2. Tap an underlined word such as **Waiting period**: *In plain words*, an example and the Hindi word.
3. Press **Close** or Esc to go back.

Get cover
Demo date and time: 18 August, 18:00

◀ SIMULATED ⓘ

New cover always starts 7 days after you ask.

Before you pay: Chhatri uses your sales data to decide claims and set your premium. After the first payment, Chhatri takes the next day's premium from your daily settlement. Chhatri reads a hospital slip when you send one, for that claim alone. You can turn each of these off in the app under Help, My data and consent. This is a prototype with simulated data.

Notice version notice-1

☐ I agree: Chhatri may use my sales data to decide claims and set my premium.
Needed for cover

☐ I agree: after the first payment, Chhatri may take the next day's premium from my daily settlement.
Needed for cover

☐ I agree: Chhatri may read a hospital slip when I send one.
Optional

Tick the two boxes marked Needed for cover to continue.

Tick the boxes marked Needed for cover to go on.

[Go to the boxes](#)

Home Claims Help

Get cover
Demo date and time: 18 August, 18:00

◀ SIMULATED ⓘ

New cover always starts 7 days after you ask.

Blocked for now. Cover starts on 25 August.

New cover starts after the waiting period

Cover starts	25 August
Premium a day	₹14.16
First payment (30 days)	₹424.80

Cover bought after an alert was issued does not pay for that alert. You can still buy cover for later.

Chhatri uses your sales data to decide claims and set your premium.

◀ SIMULATED ⓘ

SIMULATED payment link. No real money moves.

Payment link: <https://paytm.me/sim-7BFFBE>

[Simulate payment](#)

Waiting for your payment. This page updates when it arrives.

Pay the first payment to start your cover.

[Simulate payment](#)


Home Claims Help

6.19 Buy cover, with consent

Home · Get cover · scenario *buy cover*, Monday 18 Aug 18:00

Shown for Ramesh (S-0907). Cover starts only after a 7-day waiting period, so nobody can buy it once a storm is forecast. The left picture is the form, with two boxes marked *Needed for cover* and one optional. The right picture is the answer.

1. Tick the two **Needed for cover** boxes; until then the check button stays grey.
2. Press **Check price and start date**: *Blocked for now. Cover starts on 25 August. with ₹14.16 a day and a first payment of ₹424.80 for 30 days.*
3. Press **Simulate payment**. No real money moves.





6.20 Claims tracker

Tab Claims

The list of Anil's claims. One card: *Rain and lost sales*, 19 August, Paid, ₹1,380, "All steps are done. Your receipt is ready." The bar at the bottom says *Open your latest claim to see each step*.

1. Tap the claim card to open it.
2. If the list is empty it says so and tells you the next step.




6.21 Claim detail

Claims · the claim card

A **What happens next** card, then a five-step tracker. Paid and the instalment holiday carry a SIMULATED label.

1. Read **What happens next**: the credit (₹1,380 at 17:04), the lender's answer (the ₹600 instalment moves to the end of the loan, no penalty), what to do if the amount looks wrong (we reply within 24 hours) and that the cover continues (paid up to 22 August). Its lines come from the receipt and the cover, not from a model.
2. Scroll through the steps (right picture): Detected, Checked (all 9 checks passed), Decided (approved ₹1,380), Paid, Instalment holiday.
3. Use the buttons below: **See receipt, Why this amount?, This is wrong**.





6.22 Why this amount

Claim detail · Why this amount?

The sum: $\frac{1}{2} \times ₹4,380 \times 63\% = ₹1,380$ in English and in Hindi. Each number sits on a row with a **source badge** (a small label that says where the number came from and when, for example Area sales index · 17:00 · SIMULATED). Underlined words such as *Your usual day* and *Area drop* open the jargon lens.

1. Read the formula at the top.
2. Read **Your numbers**: your usual day ₹4,380, area index 37%, area drop 63%, Chhattri pays half, most paid for one day ₹2,500, and the payout ₹1,380.
3. Tap a badge to see its source.
4. At the bottom a line says what would have changed the outcome, written by the engine, not by a model. The blue bar offers **See the receipt**.





6.23 Receipt

Why this amount · See the receipt

A payout receipt that opens with **Who authorised this money**: the policy engine (code, not AI), rules `pilot-0.1`, *AI authority: None*, and **Verify this decision**, which checks the audit log. The right picture is the bottom of the receipt.

1. Press **Check the log**: *Log unbroken, N entries*.
2. Read **If you disagree**: claims officer (24 hours), insurer's grievance officer, IRDAI's Bima Bharosa portal, Insurance Ombudsman.
3. Press **Print or save as PDF**, or **This is wrong** to open the claim and send a dispute (6.27).





6.24 Slip pre-check

Help · Send your hospital slip · scenario *illness*, Thursday
21 Aug 11:20

The personal claim. Chhatri noticed Anil's shop had no sales yesterday and asked if he is well; now he sends one photo of the hospital slip. The app shows what it read, and Anil must confirm before any check runs. A line warns that the photo may go to an AI reading service (Gemini or Sarvam), so use a sample slip.

1. Press **Use the demo slip** (or **Use a blurry demo slip** to see a bad photo handed back).
2. Read the fields (Anil R. Jadhav, admitted 20 August, KEM Hospital, Parel) and the three ticks. No percentage is shown.
3. Press **Yes, this is right** to open the claim (₹1,500 approved), or **Send another photo**.





6.25 Ask Chhatri, with citations

Help · Ask Chhatri

A question box with four example questions. The answer is built from the engine's own facts and cites the policy clause it relies on. A small label under it says what wrote it: the mode (LIVE, SIMULATED or FALLBACK) and the provider, such as *rules* or *gemini*.

1. Tap an example such as **Why did I get this amount?**, or type your own and press **Send**.
2. Read the formula, *From the policy* with a clause chip (**C4.1 How much we pay**) and *Based on* with the numbers and their source badges.
3. Tap the chip to read the clause (next picture). A message that looks like a scam gets a warning banner.





6.26 The clause sheet and voice chips

Ask Chhatri · **Speak** (shown in Hindi) · flag n4_voice

The left picture is the clause sheet. The right picture shows voice, in Hindi: when the app hears an amount it asks Anil to confirm it with a chip, ₹2,500 — *is that right?*, with buttons **Right** and **Change**. **Send** stays grey until he confirms. In the static demo the voice is simulated.

1. Press **Speak**, then **Speak** again in the first-use notice (it says the voice may go to a speech service).
2. Speak, press again to stop, and check the text.
3. Tap **Right** on the amount chip, or **Change**, then press **Send**.





6.27 Grievances

Help · **Complaints and escalation** · flag

n5_grievances

Where a merchant complains. With nothing filed it says so. After filing it shows one ladder: the claims officer (you are here, we reply within 24 hours), the insurer's grievance officer, IRDAI's Bima Bharosa portal and the Insurance Ombudsman. The outside steps are labelled SIMULATED or "file it yourself".

1. Press **New complaint**, pick one of nine topics (for example **The amount of a payout**) and write a few words (up to 500 characters).
2. Press **Send**. The ladder appears with the case number.
3. Press **Mark as solved** to close it.




6.28 Consent centre and activity

Help · My data and consent · flag n6_consent

Consent is the merchant's permission to use data for one purpose. Each purpose has a switch, the date it was agreed, a **What we use** note and a **Receipt** button. The right picture is the activity log: what was used, for what and when.

1. Turn a switch off. A sheet explains the effect, for example that Chhatri stops reading slips. Confirm or keep it.
2. Press **See what was used** for the activity log; filter by Sales, Slip or Premium, or press **Check the log**.
3. If a slip is held, press **Erase this slip** to delete it on request.





6.29 Help

Tab Help

One list of everything else: what am I covered for, Ask Chhatri, complaints and escalation, my data and consent, language, and send your hospital slip. A line explains that if an amount looks wrong you open the claim and tap **This is wrong**. Below it, *Insurance words in plain language* lists the words of the jargon lens. The bar at the bottom suggests the next step: in this picture a case is open, so it says the case is with a claims officer.

1. Tap a row to open it. Back returns here.
2. Tap a word such as **Weather alert** to see it explained.




6.30 Language switch, in Hindi

Help · **Language**, or the globe icon

The Language row opens **Settings**. The app speaks Hindi, English and (with `n8_marathi`) Marathi, each in its own script, with a preview line. The right picture is Home after choosing Hindi: the same cover card, the same numbers, in Devanagari.

1. Tap **हिंदी**. A preview line shows how it will read.
2. Press **Back to Home**. The tabs now read होम, दावे, मदद.
3. Switch back with the same screen, or the globe icon.



6.31 Messages from Chhatrī: WhatsApp or Telegram

Help · **Language** · second block of Settings · flag `telegram_channel`

The second block on Settings is **Messages from Chhatrī**: WhatsApp (the default) or Telegram. It only changes where payouts, check-ins and answers are sent; cover and payouts stay the same. After choosing Telegram the screen says whether a chat is linked (here: *Not linked yet: open Chhatrī in Telegram and press Start.*), offers **Open Chhatrī in Telegram** once the bot's address is known, and, while the bot is not live, notes that *Telegram is simulated here: messages show on the console phone.*

1. Choose **Telegram** under Messages from Chhatrī.
2. With a live bot, press **Open Chhatrī in Telegram** and then **Start** in Telegram to link the chat.
3. Choose **WhatsApp** to switch back. The same switch is on the Merchant phone page (6.11).

7 The three-minute demo story

This is the stage cut of the demo, from the stage script (docs/06-delivery/stage-script.md). One person narrates and a second person clicks. Start the stack with `make demo-stage` : it runs the backend and the console with the stage flag set and synthetic data, with no auto-reload. By default it uses the AI keys in `.env` (Gemini reads the slip and answers free questions, and the footer says LIVE); `make demo-stage STAGE_AI=sim` blanks both AI keys, so the slip reader and the model are simulators and say SIMULATED. Before you start, pick **Monsoon replay**, seek to **16:40** and stay paused on the Live map. For the longer seven-minute story (illness, the mismatched slip, the blocked cover purchase, audit, policy and backtest) see docs/DEMO.md .

Time	What you say	What you do	What the screen shows
0:00-0:20 · the honest line and the hook	"Before the demo, one honest line: the policy engine, the console and the audit chain are real; shop sales, WhatsApp and the Paytm link are simulated, and every screen labels which is which. Anil runs a tea stall in Parel. In heavy rain his sales fall by more than half, and his ₹600 loan instalment is still cut. Getting paid for a day like that has typically taken 30 to 60 days. Watch what Chhatri does with no claim."	Nothing. Console parked on the Live map at 16:40, paused.	The map at 16:40, the header chip (<i>Simulated · 18</i> when nothing is live, as with <code>STAGE_AI=sim</code> and no Telegram token; with live keys it names the live services first, for example <i>Gemini · Telegram + 15 simulated</i>) and the footer line <i>Sales, alerts, KYC, payouts, lender and Soundbox are simulated</i> .
0:20-0:55 · the trigger	"A red alert from 14:00 over three wards. Zone 7: sales at 37% of expected for three hours, all 46 shops. At 17:00 the rule fires. At 17:04, with the evening settlement, 312 shops are paid. Nobody filed anything." Then: "Zone 9 also dipped, to 61%, but it had no alert. A slow day is not a loss event, so it gets nothing."	Press Play . About 12 seconds later the clock reads 17:05; press Pause .	Labels <i>Z7 · 37% · 46 shops, Z3 · 38% · 141 shops, Z12 · 47% · 125 shops</i> . The note <i>₹1,380 credited · 17:04</i> . Tiles 3 zones, 312 shops, 4 min. The Zone 7 card: 37%, 46 of 46 prepaid, ₹58,900. Z9 reads 61%, slow day.
0:55-1:25 · Anil's phone	"Anil's phone, in Hindi. 1,380 rupees with today's settlement, and the Soundbox says it out loud. No claim needed." Then: "He follows the claim: detected, checked, decided, paid. And tomorrow's 600 rupee instalment: Chhatri asks his lender for a holiday, and the lender decides. Here the lender is simulated, and it says so."	Open Merchant phone . In the app, tab Claims (दावे), then tap the claim.	The ₹1,380 card with <i>No claim needed</i> , the Soundbox line and the lender line, then the five steps, all done, with SIMULATED on Paid and on the instalment holiday.

Time	What you say	What you do	What the screen shows
1:25-1:50 · why	"His usual Tuesday is 4,380. His area fell 63%. Chhatri pays half of the lost sales. Every number says where it came from. Tap one: it names the record, its time, and that it is simulated."	Tap Why this amount? (इतने पैसे क्यों?), scroll to the area-index row, tap one source badge, read it, tap Close .	$\frac{1}{2} \times ₹4,380 \times 63\% = ₹1,380$ with a source badge on every row. The sheet says SIMULATED. A line says one more point of drop would have added about ₹22.
1:50-2:05 · Thursday, Anil falls ill	"Thursday. Anil's shop was shut all Wednesday while his area was normal. Nobody filed a claim: Chhatri checks in before he asks. He is in hospital."	On the Live map pick Anil falls ill · Thu 21 Aug , seek 11:20 , open Merchant phone and tap the chip that means <i>I am in hospital with a fever</i> (मैं अस्पताल में हूँ, बुखार है।).	The chat shows <i>Your shop has been closed since yesterday. Is everything okay?</i> and then <i>Get well soon. Please send one photo of the hospital slip</i> .
2:05-2:38 · the slip	"One photo. The app shows what it read, and asks whether that is right. If it is not sure it asks for another photo instead of guessing." After the credit: "Half his usual Wednesday is 2,150, capped at 1,500 a day. 1,500, credited with the settlement."	In the app: tab Help (मदद), Send your hospital slip, Choose from the gallery , pick <code>anil_admission_slip.png</code> (or press Use the demo slip), tap Yes, this is right . Then press Play , wait for the paid step and press Pause .	The fields Anil R. Jadhav, 20 August, KEM Hospital, Parel, three ticks and no percentage; the footer says SIMULATED (or LIVE when Gemini reads the slip). Then ₹1,500 approved: $₹4,300$ का आधा = $₹2,150$ प्रतिदिन; सीमा $₹1,500 \times 1$ दिन = $₹1,500$, and the paid step done and SIMULATED.
2:38-2:52 · Ask	"He can ask in his own words. The answer cites the policy clause, and the footer says what answered. Here it says rules: the policy engine, not a language model. With a live key, free questions go to the model and the footer says live and names it."	Tab Help, Ask Chhatri , type <i>Why did I get this much?</i> in Hindi (मुझे इतने पैसे क्यों मिले?) and press Send .	An answer with a rupee figure, a clause chip (<i>C4.1</i> and its neighbours), source badges, and a footer that reads LIVE and rules.
2:52-3:00 · the audit chain and the close	"Every step you just saw is in a hash-chained log. The chain reads valid. The AI builds the case; code decides the money. Sales, alerts, KYC, payouts, the lender, WhatsApp and the Paytm link are simulated and labelled on screen."	Open Audit and press Verify chain .	<i>Chain valid · N entries</i> and the footer disclosure line.

IF THE CLOCK RUNS LATE

Drop items in this order: the Zone 9 line, the source-badge tap, the Ask step, the ₹1,500 credit step (say it instead). **Never drop the slip pre-check, and never the closing disclosure sentence.**

When a judge asks (not in the three minutes)

Question	Open	Say
"Who approved this money?"	Anil's app: Claims , the claim, See the receipt (6.23)	The receipt opens with who authorised the money: the policy engine under rules <code>pilot-0.1</code> , AI authority none, and Verify this decision , which checks the audit log.
"What happens to Anil after the payout?"	Anil's app: the paid claim (6.21)	The What happens next card: the credit, the lender's answer (SIMULATED), how to dispute and when he hears back, and the date the cover is paid up to. Every line comes from the receipt and the cover.
"How do you know the AI is any good?"	The Evals page (6.16). It is not in the stage flag set: use <code>make dev</code> with every flag, or say it.	One stored run on synthetic questions we wrote: 50 Ask questions run live, no answer with a forbidden statement, 30 of 31 citing the right clause, 2 of 50 with a rupee figure not traced to a fact; slip reading and voice not measured. Quote the page, with its n.
"How did you price it? What if the price is wrong?"	Backtest , scroll to Price the cover (6.14)	Every zone is priced from the backtest's own triggers. Z7 is ₹18.62 a day at the published rules; press 55% and the premium, the loss ratio at today's price and the false payouts move; drag the share and the premium and the loss ratio move. "Planning figures, not an actuarial price; the pilot sets the price."
"Is the build real?"	A terminal: <code>make judge</code>	One line per check, ending <code>CHHATRI JUDGE READY ✓</code> : the data files against their fingerprints, 70 demo checks, the audit chain after the payout.

Spare 15 seconds: on the Live map after 17:00, click **Z9**, then **What if...**, drag the three hours to 49 and tap **Rain**. The panel reads *Would fire: yes, 51% drop*, and nothing is saved (6.7). Spare 30 seconds, the live tests of the Policy page (6.15): load **Illness mismatch** for HUMAN (a name that does not match is never decided by the AI) or **Buy cover** at 18:10 for BLOCKED (cover bought after an alert is told the truth, never approved).

Golden numbers to check against the screen

Fact	Value
Replay day	Tuesday 19 August 2025, 08:00 to 20:00 (simulated)
Alert	A-20250818-01, red rain, zones Z3, Z7, Z12, issued Monday 18 Aug 17:30, valid 14:00 to 20:00
Indices at 17:00	Z7 37% · Z3 38% · Z12 47% · Z9 61% (64 shops, slow day, no alert)
Timeline	Decisions 17:00 · credits, WhatsApp and Soundbox 17:04 · lender asked, instalments paused 17:05
Totals	3 zones triggered · 312 shops paid · 4 min trigger to money · Z7 total ₹58,900
Anil	Usual Tuesday ₹4,380 · payout ₹1,380 · instalment ₹600 paused by his lender
Personal claim	Usual Wednesday ₹4,300 · half is ₹2,150, capped at ₹1,500 for 1 day · credit 4 minutes after the decision
Cover link (Ramesh, S-0907, Z3)	₹424.80 for 30 days (₹14.16 a day). Cover bought on Monday 18 Aug is blocked and starts 25 August
First case	C-2291 after a fresh load of a scenario

If the console ever shows something different, reload the scenario from the list, and if it still differs, run `make demo-check` before going on.

8 Features list

Each feature has an id. **N** is a new feature, **X** is a fix to something that already existed, **H** is an idea adopted from another project. "Flag" is the switch from section 5. "none" means it is always on.

N: new features

Id	In plain words	Flag
N1	The merchant mini-app: a phone-sized app with home, coverage, buying cover, claims, receipt and help.	n1_miniapp
N2	Ask Chhatrī: ask a question in your own words and get an answer that cites the policy.	n2_ask_chhatrī
N3	Slip reading with a pre-check: the app shows what it read and you confirm before anything is checked.	n3_slip_precheck
N4	Real Hindi voice: speak a question and hear an answer.	n4_voice
N5	Complaint ladder with response clocks: who to complain to next, and by when.	n5_grievances
N6	Consent centre: see, turn off and log each permission; erase your slip.	n6_consent
N7	A static build of the console that runs with no backend. The build exists; it has not been put on a public address.	none
N8	Marathi as a third language (a draft, text only).	n8_marathi

X: fixes

Id	In plain words	Flag
X1	The frontend tests pass, every time, even on a busy laptop.	none
X2	A bad claim is never stored: the "usual day" amount must be a properly rounded published value.	none
X3	A missing zone price now fails loudly instead of quietly charging the ₹2 minimum.	none
X4	The lender decides the instalment holiday. Chhatrī only asks. Before this fix the pause was unconditional.	x4_lender_request
X5	Every route answers an unknown merchant id with a clean "not found".	none
X6	Provider panel with a fallback switch: see which services are live and force a fallback.	x6_provider_panel
X7	An honest-wording test: no screen claims a simulated thing is live.	none
X8	No loan offers to a merchant in distress, and a cap on messages.	x8_distress_guard

H: ideas adopted from other projects

Id	In plain words	Flag
H1	Claim tracker: Detected, Checked, Decided, Paid, Instalment holiday.	n1_miniapp
H2	"Why this amount" with source badges.	none (view in n1_miniapp)

Id	In plain words	Flag
H3	A payout receipt you can print.	none
H4	Honest-wording test (same as X7).	none
H5	The merchant confirms what was read from the slip.	n3_slip_precheck
H6	A public demo with no login (the static build; not published yet).	none
H7	A provider panel that shows how the AI services are doing, with a fallback switch.	x6_provider_panel
H8	Real operations counts taken from the database, never projections.	h8_ops_strip
H9	No loan offers in distress, plus a message cap.	x8_distress_guard
H10	A deterministic core: the AI has no authority over money.	none
H11	Bilingual message templates instead of model translation.	none (Marathi: n8_marathi)
H12	Visible, measured test counts in the README.	none
H13	Source badges: every rule, number and clause says where it came from, and when.	none
H14	A counterfactual in every explanation: what would have changed the outcome, written by the engine.	none
H15	Slip pre-check: document type, checklist, and "ask, do not assume" when the reading is unsure.	n3_slip_precheck
H16	Defence against orders hidden in slips and chat.	none
H17	Ask Chhatri cites policy clauses; every number comes from the engine.	n2_ask_chhatri
H18	Voice confirmation chips for amounts and dates.	n4_voice
H19	Scam-message warning in chat (OTP requests, advance fees, "guaranteed claim", urgency).	n2_ask_chhatri
H20	Jargon lens: tap an insurance word for a plain explanation.	n1_miniapp
H21	Next-step bar: no screen or reply is a dead end.	n1_miniapp
H22	Grievance ladder with response clocks and a "who owns this complaint" router.	n5_grievances
H23	Consent activity log and "forget my slip".	n6_consent
H24	What-if drawer for judges (read-only), and the pricing simulator.	h24_whatif
H25	Published AI evaluation: a number is shown only once it is measured.	h25_evals
H26	Every AI reply carries its mode (live, simulated, fallback), provider and reason.	none

One later addition has no id: the Telegram chat app (flag `telegram_channel`), a second place for Chhatri's messages next to WhatsApp. Console-only flags: `console_polish` (Presenter mode, the moment card and projector type), `h8_ops_strip` , `h24_whatif` , `h25_evals` and `x6_provider_panel` . All 15 flags are listed with their effect in section 5.

9 Troubleshooting, FAQ and glossary

Troubleshooting

Symptom	What to do
Numbers differ from this manual	Reload the scenario from the list at the top left (it is deterministic). If they still differ, run <code>make demo-check</code> .
A page or button is missing	Its flag is off. Set the flag in both <code>CHHATRI_FEATURES</code> and <code>VITE_FEATURES</code> (section 5) and restart. A flagged route answers 404 when off.
Page is blank or reset after pressing reload in the static mock	The mock backend lives in the page, so a reload restarts it. Load the scenario again and use the menu links, not reload.
<i>Reconnecting pill</i>	Wait. The stream resumes by itself. If it lasts over 10 seconds, reload the page; the state is kept by the backend.
No sound	Press Enable sound in the header once. Every voice bubble also has a play button.
Phone shows nothing after a slip or an officer approval	The clock is paused and money arrives 4 simulated minutes later. Press Play , or press +15 min .
Replay stuck	Pause, press +15 min , or seek to the time you need.
Map has no street tiles	Nothing to fix. Without internet, or when the map server does not answer, the console draws the ward outlines on a plain background.
<code>make demo-check</code> fails many steps	Check that <code>n3_slip_precheck</code> and <code>n6_consent</code> are not on, and that <code>x4_lender_request</code> is.
After saving a backend file, the replay and cases vanish	<code>make dev</code> restarts the backend when a file changes, which clears everything held in memory. Reload the scenario and seek again.
Console does not load with <code>make up</code>	Run <code>docker compose ps</code> and <code>docker compose logs backend</code> . As a fallback run <code>make dev</code> (console on port 5173).
Clock holds at 17:00 for 30-60 seconds (n8n live)	Expected: 312 payout runs go through n8n one after another. To avoid it, set <code>CHHATRI_STACK_N8N_URL=</code> (empty) in <code>.env</code> .
<code>/api/preflight</code> shows an item that is not ok	An artefact is missing. Run <code>make data</code> (slow) or restore the committed files.
An AI answer reads SIMULATED or FALLBACK although a key is set	Run <code>make check-keys</code> . A Gemini key needs a model id (<code>GEMINI_MODEL</code>) and <code>CHHATRI_DATA_IS_SYNTHETIC=true</code> . On the free tier each Gemini model allows only about 20 requests a day, so a spent model hands over to <code>GEMINI_BACKUP_MODELS</code> , then to Sarvam, then to a fixed template. The label and the provider panel say which link answered and why.
Telegram reads SIMULATED, or no chat linked yet	A live bot needs the flag <code>telegram_channel</code> , <code>TELEGRAM_BOT_TOKEN</code> and <code>CHHATRI_DATA_IS_SYNTHETIC=true</code> . Then open <code>https://t.me/<bot name>?start=S-0142</code> on a phone and press Start. Only one machine may poll a token: run every other one with <code>TELEGRAM_POLLING=false</code> .

Symptom	What to do
The pricing section says it needs the backend's pricing table	You are on the static mock, which has no pricing table. Run <code>make dev</code> with <code>h24_whatif</code> on (section 5) to see Price the cover .

Questions a newcomer asks

Is Chhatri a real insurance product?

No. It is a hackathon prototype. No insurer, lender or Paytm team has agreed to take part, and nothing here is an approved product.

Does real money move?

No. Payouts, the lender and the payment link are simulated and labelled.

Who decides the money, the AI or the code?

The code. A policy engine reads `rules.yaml` and decides. The AI only builds the case and explains.

What happens when the AI is unsure or unavailable?

A doubtful slip goes to a claims officer, never to an automatic payout. If an AI service is not available, a simple template or simulator answers and the label says so.

How is the payout worked out?

For a storm, half of the lost sales: $\frac{1}{2} \times \text{the usual day} \times \text{the drop}$. For Anil, $\frac{1}{2} \times ₹4,380 \times 63\% = ₹1,380$. For illness, half of the usual day, capped at ₹1,500 a day.

Why did Zone 9 get nothing?

Its sales fell to 61% on a day with no weather alert. That is a slow day, not a loss event.

Does Chhatri pause a loan?

No. It asks the lender, under a rule agreed in advance, and the lender decides. In the demo the simulated lender says yes.

Why can a merchant not buy cover when a storm is forecast?

Cover starts after a 7-day waiting period, with a 72-hour look-ahead at alerts. Otherwise people would buy cover only on the eve of a storm.

Do I need API keys or an internet connection?

No. Without keys everything is simulated, and the static mock runs in the browser with no backend. The map pictures come from the internet; without it the console draws the ward outlines instead.

Which languages does it speak?

Hindi and English. Marathi is a draft, in text only, behind `n8_marathi`. Voice is Hindi and English.

Is any real person's data used?

No. Every merchant, slip and sales figure is synthetic. AI services are called only when that is declared.

How can I check that nobody changed the log?

Open the Audit page and press **Verify chain**, or press **Check the log** on a receipt in the mini-app.

Which chat apps does it use?

WhatsApp is the default and is simulated here. With the flag `telegram_channel`, a merchant's messages can go to Telegram instead, and the Telegram bot can run live. Only where messages go changes: cover and payouts stay the same.

How do you know the AI is any good?

Open the Evals page. It shows one stored run on synthetic questions the team wrote, with its misses, and it says NOT MEASURED for slip reading and voice. No person has graded an answer yet.

How is the cover priced?

Each zone's premium is worked out from the backtest's own triggers, with a 35% loading and a minimum of ₹2 a day. For Anil's zone it is ₹18.62 a day. On the Backtest page, **Price the cover** lets you move the levers and see the effect. These are planning figures on simulated sales, not an actuarial price; the real price is the open question for a pilot.

Glossary

Term	Meaning
Alert	A weather warning from the India Meteorological Department with colours: Green, Yellow, Orange, Red. Chhatri looks at Red and Orange.
Audit log	A record of every step. Each entry holds the fingerprint of the one before, so a change is detected.
Backtest	A replay of past monsoons to see what the rules would have paid. Here it uses simulated sales and real rainfall.
Basis risk	The gap between what the trigger says and what the merchant really lost.
Case	An open dispute or doubtful claim, with a 24-hour answer clock, such as C-2291.
Claim	A request to be paid after a loss.
Consent	A permission to use your data for one named purpose. You can withdraw it.
Counterfactual	A line that says what would have changed the outcome.
EDI	Equated daily instalment: the daily loan repayment taken from a merchant's settlement.
EDI holiday	A pause of one instalment, which moves to the end of the loan. Chhatri asks; the lender decides.
Expected sales	What a shop normally sells at that hour or on that weekday, estimated by a model.
Feature flag	A switch that is off until you turn it on.
Hospital cash	A fixed daily amount for each day in hospital, whatever the bill.
Index floor	The share of expected sales that an area's sales must fall below to count as a loss. In the published rules it is 50%.
IRDAI	The Insurance Regulatory and Development Authority of India. Its Bima Bharosa portal takes insurance complaints.
Jargon lens	Tap an insurance word to see a plain explanation.
KYC	Know Your Customer: the identity record of the merchant. A slip's name must match it.
Lender	The business that lent the merchant money. It decides about an EDI holiday.

Term	Meaning
Loading	The extra amount added to the expected payouts when a premium is priced. In the published rules it is 35% of the premium.
Loss ratio	Payouts divided by premiums. Above 100% means a zone pays out more than it collects.
Parametric	Insurance that pays when a set trigger happens, without measuring each shop's loss.
Payout	Money paid after a claim is approved.
Policy engine	The code that decides APPROVED, REFERRED or DECLINED. Only it can approve money.
REFERRED	A decision that sends the claim to a person.
Replay	A run of a past day, with a clock you control.
rules.yaml	The file that holds the policy rules and numbers.
Sales index	Actual sales divided by expected sales for a zone.
Settlement	The daily payment from Paytm to a merchant, after fees and loan deductions.
Slip	The hospital paper (admission or discharge) used as proof for a hospital claim.
Soundbox	The speaker on a merchant's till that announces payments.
Source badge	A small label that says where a number came from and when.
Telegram bot	Chhatri's bot on the Telegram chat app, a second place for a merchant's messages next to WhatsApp. It is simulated unless a bot token is set.
Trigger	The event that starts a claim.
Waiting period	The days after buying cover before it can pay. Here, 7 days.
Zone	A group of neighbouring shops in one area. The demo uses Z3, Z7, Z9 and Z12.

WHERE TO READ MORE

The documents index is `docs/README.md`. The demo script is `docs/DEMO.md`, and the three-minute stage script is `docs/06-delivery/stage-script.md`. How the parts fit is in `docs/ARCHITECTURE.md`, and every outside service is described in `docs/INTEGRATIONS.md`. Data credits: map tiles © OpenStreetMap contributors (ODbL); ward boundaries from DataMeet Mumbai BMC wards (CC BY-SA 2.5 India); weather data by Open-Meteo.com (CC BY 4.0).